



- Text data – Unicode/UTF-8 quick recap
- Special text formats: JSON, HTML, XML
- JSON examples: Twitter
- Reading files line by line
 - Basic file operations
- Retrieving data from the web
 - This will usually be HTML for pre-existing tables, JSON or XML for “dynamic” data
- Web scraping examples
 - Free-form vs. via REST API



NAVAL
POSTGRADUATE
SCHOOL

OS4118

Statistical and Machine Learning

Text Data

Prof. Sam Buttrey

Fall AY 2020

The Nation's Premiere Defense Research University

Monterey, California
WWW.NPS.EDU



- **Unicode** is a scheme to represent characters from all the world's languages
- Can display about 1M characters (“code points”), broken into 17 “planes”
- The characters in plane 0, the “basic multilingual plane,” have two-byte addresses 0x0000 to 0xFFFF
 - Superset of ASCII, which is 0x00 to 0x7F



- The **UTF-8** encoding of Unicode is byte-ordered and is the standard for web pages, XML, and, really, everything else
 - Except that Microsoft's "Save as Unicode" uses the different UTF-16
- ASCII require 1 byte (backwards compatibility); others require 2-6
- Example: Euro sign € is \U20ac
- Represented in UTF-8 as hex e2 a2 8c



The NUL Character

- NUL (0x00) requires special thought
- For years it served as a de facto “end of string” character because it serves this role in C programs (see ?readBin)
 - Issues: learning the string length requires traversing the whole string; programmers often forgot to allocate the extra character; strings and binary must be treated differently...



- UTF-8 encodes NUL as 0x00, but...
- Lots of legacy code uses UTF-8 + NUL-terminated strings
 - “Modified UTF-8” treats NUL in a special way, encoding it as 0xC080, so that no single UTF-8 byte is ever 0x00
 - This is not permitted by the UTF-8 standard
 - NUL characters are generally not permitted in R, so if they’re in your file, watch out
 - Moral: Be aware of NUL issues



- **Delimited** (e.g. CSV, tab-separated)
 - Flexible
 - The separator itself takes up room
 - The separator can be mistaken for legitimate contents – commas are common in free text
 - Europeans more often use ;
- **Fixed-format**
 - Sometimes easier to handle, smaller...
 - Although format specs can be a pain to use



```
Name , ID , State , Amount1 , Date1  
Buttrey , 003682X , CA , 102.45 , 11/09/2016  
Jenkins , 001926Z , WA , 2130.40 , 11/30/2016  
Lee , 000411R , MI , -1.36 , 11/18/2016  
Rimsky-Korsakov , 6076T , QB , 642.28 , 11/28/2016
```

- Generally carries column headers, but you'll still need documentation
- Again, the comma is a common choice for delimiter but causes trouble with free-form text (even “comma-free” entries like city names!)



Data:

```
Buttrey.....003682XCA 0001024520161109
Jenkins.....001926ZWA 0021304020161130
Lee.....000411RMI-0000013620161118
Rimsky-Korsa006076TQB 0006422820161128
```

Format (shown here COBOL style):

Name	PIC X(12)
ID	PIC X(07)
Amt1	PIC S9(06)V99
Dt1	PIC 9(08)



- HyperText Markup Language (**HTML**) is a scheme for telling browsers how to display pages
- Text together with “markup” tags
 - `<html><head>...</head>`
 `<body>Text goes`
 `here</body></html>`
 - “Head” includes title, metadata like author
- Often tags come in pairs: `bold`,
`red`



- Markup gives display information to the browser, which is in charge of “rendering” (displaying) the page
- The set of tags is fixed by standard, although inevitably browsers disagree
- HTML supports a list of special characters; usually our tools interpret these properly
- HTML supports UTF-8, which will be the default in the upcoming **HTML5**
- One important tag is `<table>`



- Enclosed in `<table></table>`
- Rows in `<tr></tr>`, entries in `<td></td>`
 - `colspan`, `rowspan` to span multiple col/rows
 - HTML tags are not case-sensitive
 - Table rows can span many lines of text file
- Header in row 1 or inside `<th></th>`
- Tables often nested for formatting reasons

- `readHTMLTable()` in R library XML or the newer library `htmltab` will try to extract table data...
 - ...producing a list of tables found...
- ...but HTML is often non-standard; you may have to acquire, decode it yourself
 - With, e.g., `GET {httr}`, `getURLContent {RCurl}`
- Examples: EIA tables



- **XML** is a standard mechanism for storing and transporting (but not displaying) data
- It looks a little like HTML, but HTML's focus is on the display
- XML is very Unicode-friendly, usually using UTF-8
- XML is text-based, so it's not particularly well-suited for floating-point data



- **XML**

- Very structured, easier to impose content restrictions at creation time (properly-formatted dates, factor levels...)
- Much bigger than just the data
 - Zillions of tags take up most of the room
- Not friendly to binary data
- Special editors helpful
- We will never write a parser for XML; there are libraries available for every serious language

- HTML comes with a fixed, defined set of tags (<p>, <i>, , <table> etc.)
- In XML, we define our own (case-sensitive) tags for the purpose at hand
- So XML isn't quite a language; it's a set of tools for defining a new language
- E.g.: one such language, called **KML**, is what is used by Google Earth
 - Tags like Placemark, Polygon, coordinates...



```
<customer><Name>Buttrey</Name>  
<ID>003682X</ID><State>CA</State>  
<Amount1>102.45</Amount1>  
<Date1>20161109</Date1> ... </customer>
```

```
<customer><Name>Jenkins</Name>  
<ID>001926Z</ID><State>WA</State>  
<Amount1>2130.40</Amount1>  
<Date1>20161130</Date1>...</customer>
```



- For NAVAIR Data Challenge 1, participants were given XML files with aircraft maintenance data
- Saved “as Unicode” in Windows, which means UTF-16, not UTF-8
 - Every ASCII character – so, every character in the file – took up two bytes instead of one
- No “new-line” characters were present
- It’s important that producers and consumers of data be on the same page!



- XML library at CRAN
 - Includes our `readHTMLTable()` function
- Functions to ingest, deparse, store, and write XML
- **XPath** and **XQuery** are query languages to select nodes from XML trees
 - Also can extract and compute on values
 - `Xpath {XML}`



- **JSON** (“JavaScript Object Notation”)
- Began as a mechanism for server-browser communication
 - Every browser understands JavaScript
- Now used as a storage format
 - E.g. Twitter data
- Braces surround name-value pairs, separated by commas

- `{ "Name" : "Buttrey" ,
 "ID" : "003682X" , "State" : "CA" ,
 "Amount" : 102.45 ,
 "Date" : "2016-11-09" }`
- Special values `true`, `false`, `null` and numbers don't need quotes
- Can contain nested arrays (in square brackets) or objects (in braces)
- Always best handled with a library



- Free-form
 - Free-form text is never your friend
 - Inconsistencies in format, spelling and typography; end-of-line hyphens, all cause problems
 - Extracting meaning is hard

Buttrey, a Californian, sent his payment on November ninth of last year, in the amount of \$102.45...



Some Niceties of Text in R

- Individual string lengths given by `nchar()` (compare `length()`)
- How many chars in the Euro symbol?
 - A: One, but it takes three bytes
 - Compare `nchar ("\t\t\n") = 3`
 - And `nchar ("\Ud7a1", type="width")`
- We often see the empty string `" "`, which has length 0 – but it's hard to distinguish that from strings like `" "`



- Missing character values look like NA
 - This is not quite the same as numeric NA
 - Definitely different from characters "NA"
 - `nchar(NA) = NA` by default*
- Strings can be unreadably long
 - `strwrap()` formats paragraphs
- There are lots of tools for common tasks like tokenization, “lemmatization,” removal of stopwords...



- Lots of R text is stored as **factors**, which are internally `<integer + labels>`
- Factor strengths: persistent levels – but this is also their weakness
- Factors are more suitable for R “long vectors” (length $\geq 2^{31}$) than characters
 - Otherwise, I advise you to avoid factors
- “Never” use `as.numeric()` on a factor
 - Use `as.character(f)` or `levels(f)[f]`



- `substring(vec, 4, 7)` produces a new vector, same length as `vec`, with characters 4-7 from each element of `vec`
 - Can also be used to assign (but not extend)
 - Start and stop arguments are vectorized; this allows one command to do a replacement for every element in a vector
 - Or, e.g. `substring(a, 1, nchar(a)-3)` # truncate 3 chars from each
 - See also `startsWith()`, `endsWith()`



- `paste()` combines arguments in a vectorized way
- Within one element, separator is `sep`, defaults to space (`sep=" "`, `paste0()`)
- Collapse into one string: `collapse=""`
 - I often use `cat (paste (vec, collapse = "\\n"), "\\n", file = outfile)`
 - Compare `write.table()`
 - `file="clipboard"/"clipboard-128"` on Windows or `pipe("pbcopy", "w")` and `pipe("pbpaste")` for Mac clipboard



Tools for Text in R (cont'd)

- `sprintf ()` formats data as text
 - Including leading/trailing zeros, scientific notation, hex, strings
 - Vectorized
- `format ()` is handy for lining things up in reports and, particularly, for dates
- R has a big set of tools that use **regular expressions** to search, split, replace
 - (The subject of a separate lecture)