



NAVAL
POSTGRADUATE
SCHOOL

OS4118

Statistical and Machine Learning

Regular Expressions

Prof. Sam Buttrey

Fall AY 2020

The Nation's Premiere Defense Research University

Monterey, California

WWW.NPS.EDU



- Common in text-matching applications: every scripting language handles these
- A “regular expression” (regex) describes a set of rules to match or not match strings
- Generalizes the “wildcard” match (which is called a “glob” – see `glob2rx()`)
- Can be very powerful and very complicated



“Some people, when confronted with a problem, think ‘I know, I'll use regular expressions.’

“Now they have two problems.”

-- Jamie Zawinski



- Regular expressions come in a number of varieties, each with slightly different rules
- Two of the main types are **POSIX**
 - ...which includes “basic” and “extended”
 - Often with “GNU extensions,” as in R
- And **PCRE**, sometimes called “Perl,”
 - ...although the Perl language uses its own variety which is a little different from PCRE
- Python re is “similar to Perl”



- In R, `grep` (and `regexpr` and relatives) use extended POSIX/Gnu by default
- `perl = TRUE` selects PCRE
- I've been using POSIX, but PCRE are also very popular
- Be sure you know which kind you're using as you look for help etc.

- `R grep ()` matches strings in a character **vector**
 - By default, returns indices of matching ones
- Common implementation `egrep` in the bash shell matches **lines** in text file(s)
 - By default, returns every line that matches
- Lots of options to modify behavior, but regular expressions are essentially for finding **text** within **lines**



- The goal is to identify strings that match a **pattern**
- Letters, numbers, and the space character match themselves in a pattern:
 - `grep ("cat", vec) # lines w/ "cat"`
 - `grep (" CA", vec) # find " CA "`
 - Don't include spaces "for readability"
- If the pattern is fixed, we can use `fixed=TRUE` in R for speed, convenience

- Some characters have special meanings in patterns

- These are all non-alphanumeric and will usually need to be **escaped** in order to have literal meaning, depending on the variety of regex

- E.g. \$ means “end of line”

`grep ("$" , vec) # find lines that end (!)`

`grep ("t$" , vec) # find lines that end
with a lower-case t`

- Special characters:
 . \ | () [{ ^ \$ * + ?
- To use one of these as itself, “escape” it by preceding it with a backslash...
- ..except that backslash is already a special character in R; we type it as \\

```
grep ( '\\$', vec ) #lines w/$
```

```
grep ( '\\\\', vec ) #lines w/ \
```

```
= grep ( '\\', vec, fixed= TRUE )
```

- `^` matches start of line, `$` matches end
 - `grep ('^The', vec) #lines starting with The`
 - `grep ('^The', vec, ignore.case = TRUE) # ignore case`
 - `.` (dot) matches any one character, so `grep ('^.$', vec)` matches lines with exactly one character (dot is special)
 - `grep ('^\\.$', vec)` matches lines with exactly one dot and nothing else



More grep Arguments

- `grep ()` tells you the indices of the strings that match (as a numeric vector)
- With `value = TRUE`, it returns the matching strings themselves
- With `value = FALSE` and `invert = FALSE`, gives the indices of strings where there is **no** match
- Also useful: `ignore.case` (default FALSE)
- `grep1 ()` returns a logical vector with TRUE for match and FALSE for no match

- [] introduces a **class** of characters
 - `grep ('ae', ...)` matches lines with `ae`
 - `grep ('[ae]', ...)` matches lines with `a` or `e`
- Ranges are permitted:
 - `grep ('[0-9]', ...)` matches lines with a digit, but beware collation sequences for letters
- A number of classes are predefined:
 - `'[:alpha:]'` matches a letter;
`[:upper:]`, `[:digit:]`,
`[:alnum:]`, `[:punct:]`, and more
are available



- A character class starting with **^** **excludes** characters in the class
 - `grep ('[^t]...' # lines w/ at least one non-t`
 - **^** elsewhere than at the start refers to itself, so `egrep '[t^]'` gets lines with t or **^**
- Classes are for single characters: use **pipe** for “word1 or word2”:
 - `egrep 'NA|[Mm]issing'` finds NA, missing, Missing – **parentheses** establish precedence



- `\w` matches a character “in a word” – that is, letter, digit, underscore
 - `= [[:alpha:]]` ; `\W = [^[:alpha:]]`
- `\s`, `\S = [[:space:]]`, `[^[:space:]]`
- `\b` matches empty string at beginning or end of word; `\<` and `\>` match at start, end; `\B` matches empty inside a word
 - 'dirty `\Brat`' doesn't match 'dirty rat'



- ? matches 0 or 1 times (“item is optional”)
- * matches 0 or more; + 1 or more times, so...
- '^ [^t] * \$' matches lines with no t
- '[0-9] ? [[:alpha]] +' matches lines w/
0+ digits, space, then a bunch of letters
- {m,n} matches m-n times, so...
- '[0-9] { 2 , 4 } ' : 2-4 digits (or more, right?)
- '\\< [0-9] { 3 } '\\>' : **word** w/exactly 3 digits



- Suppose an e-mail address is “valid” if:
 - “Local part” starts with a letter or digit, and subsequent characters are letter, digit, dot;
 - An @ separates the local and domain part
 - The domain part is made of two or three “words” of letters and/or digits, separated by dot(s)
 - (The actual rules are mostly more permissive, though double-dots forbidden)
- Find lines with (only) e-mail addresses



- Start with alpha, follow with zero or more characters from (alpha or dot)
– $[[\text{:alnum:}]]([\text{:alnum:}]|\backslash.)*$
- Then the @ sign, representing itself: @
- Then one or two (word, dot) pairs, optionally followed by one more word:
 $([\text{:alnum:}]]+\backslash.)\{1,2\}[\text{:alnum:}]]+$



R Relatives, part 1 (recap)

- `grep()` tells you the indices of the strings that match (as a numeric vector)
- With `value = TRUE`, it returns the matching strings themselves
- `grepl()` returns a logical vector
- In each case the function describes an attribute of the containing string, not so much about the match itself
 - Does the string match the pattern or not?

- Tools for extracting the matched text
- `regexpr ()` returns a vector telling **where** in the string the 1st match occurs (or `-1`), plus an indication of match length
- Of course lengths can vary, depending on the type of pattern being matched
- `gregexpr ()` describes all matches
 - Use these with `regmatches ()` to extract the matching part

- By default matching is maximal, so the match length of 'b.*a' in '**b**anana pe**a**r' is 9
- Adding ? to the repetition operator will make the match minimal
 - `[[:alpha:]]+` – get the maximal set of words and spaces, starting at character 1
 - `[[:alpha:]]+?` –first word/space char



Relatives 3: Search and Replace

- R commands `sub()` and `gsub()` let you replace strings with other strings
 - `sub()`: first match; `gsub()`, all matches
- Stuff that matches inside parentheses can be used in the replace; those are **backreferences**, referred to by `\1`, `\2`, ...
- Ex: turn “first last” names into “last, first”
- ```
sub ("([[:alpha:]]+) ([[:alpha:]]|)+",
 "\\2, \\1", nm)
```



- Think of regex as a language, of which we now know a useful subset
- Features include conditionals, lookaheads/behinds/around, others
- Most flavors handle Unicode/UTF-8, but support differs from place to place
- Expect slight variations from one implementation to the next